



Claude users don't struggle to find features. They struggle to know when to use them.

80% of Claude power users feel frustrated repeating context for the same tasks. Yet only 12% use Skills regularly.
Survey n=25 · Apr 2026.

The Pain

72% repeat work tasks weekly
80% re-explain context often (Q5 ≥ 3)

"It's mostly frustrating. Using same prompt again and again is frustrating and time consuming."
— Divyansh, Video Editor

The Workaround

56% copy-paste from old chats
40% retype manually · 20% save in Word/Notes

"I make a Word document or a text file where I save my prompts."
— Kishan, Project Manager

The Paradox

88% don't use Skills regularly
76% fear a Skill performs worse

"I haven't used the Skill feature... I don't think I require it for now."
— Vijay, Cybercrime Analyst

Why this gap matters now

Pattern density

76% repeat structured work weekly.
Strongest signal AI can detect today.

Cost waste

Power users on Pro/Teams burn tokens re-prompting context Claude already saw.

Native distribution

Claude has ~26M MAU.
Decision layer ships natively, no separate install.

Untapped capability

6+ capabilities shipped
— Skills, Agents, Workflows. Adoption gap, not feature gap.

Not every repeat should be a Skill. We tell you which ones should.



26M Claude users. Eight tools chasing them. Zero solve the decision

~26 M

Claude consumer MAU · Apr 2026
18.9M web + 7.38M mobile

70%

of Fortune 100 use Claude;
Enterprise rollout signals

\$ 14B

Anthropic run-rate · ~\$2.5B Claude
Code alone. 6+ capabilities shipped

Backlinko, Demandsage, Anthropic public reporting · Apr 2026

TAM ~ 26 M All Claude consumer MAU

SAM ~ 14.4 million
Power Repeaters with recurring task
friction (56% behavioral filter)

SOM Y1 ~ 720K Native-feature early
adopters (5% benchmark from Notion
AI, Slack AI launches)

Behavioral filter, not demographic. Sized using Anthropic
public data + survey-derived behavioral filters (n=25, Apr 2026).

*THE GAP - 8 tools mapped across 4 categories. None
detect repetition. None route to the right capability.
None preview cost-quality before commit. The market
builds execution tools. The decision layer is empty.*

*Real competition isn't ChatGPT. It's a prompts.docx on
someone's desktop. SkillSense replaces both*

Option	Helps user do what?	What user still has to figure out	Why SkillSense wins
Claude / ChatGPT	Execute tasks with AI	Whether this repeated task should become a Skill, GPT, or workflow	SkillSense identifies the pattern before the user notices it
Notion / Docs / saved prompts (used by 20% of our users)	Store reusable formats	Whether the task is reusable enough to formalize	SkillSense turns hidden repetition into an explicit recommendation
Zapier / LangChain / CrewAI	Build automations and orchestrations	Which capability to use, and when the task is worth converting at all	SkillSense sits earlier in the journey — it helps users decide, not just build
 SkillSense	Detect + recommend + nudge	Almost nothing beyond confirmation	Fills the missing decision layer in the market

*We are not another workflow builder.
We are the decision layer before workflow creation.*

Power Repeaters drive the wedge. Reluctant Adopters tell us where to draw the line.

The problem is most acute for users who repeat structured work every week.

Casual Explorers

Monthly or less
Exploratory users
No stable patterns · low frustration
24% of survey

Power Repeaters

Daily / several× weekly
PMs, content, analysts
High pattern density · strongest friction
76% of survey · PRIMARY



Reluctant Adopters

Daily, resist templating
Senior devs, designers
Output uniqueness matters
Validation segment

Heavy Devs

Claude Code / API heavy
Engineering workflows
Advanced orchestration needs
Out of V1 scope


Why these two segments? Why Power Repeaters AND Reluctant Adopters?

- 🎯 Behavioral signal density — 76% repeat work weekly. strongest pattern signal for AI detection
- ⚡ Highest-stakes pain — 100% of Power Repeaters accept the time tradeoff · casual users don't
- 🚀 Easiest to reach — already inside Claude · no separate install, no behavior change
- ✅ Trust validation — if Aniruddha's skepticism is earned, everyone else's adoption follows
Reluctant Adopters aren't the wedge — they're the proof.



Kishan, 32
Project Manager,
Hyderabad

Core sentiment: *Mujhe bas pattern dikha do — main save kar dunga.*
Behavior: Repeats weekly docs and already saves prompts outside Claude.
Core need: “Just detect the repeat for me and let me save it fast.”



Aniruddha, 28
Full-Stack Developer,
Bhopal

Core sentiment: *Templating sab kuch flatten kar deta hai — har project ka apna flavour hota hai.*
Behavior: Uses Claude daily, but resists rigid reuse when context changes.
Core need: “Show me the saved version preserves variation before I commit.”


The Power Repeater wins us the wedge. The Reluctant Adopter tells us when not to nudge.

Jobs To Be Done (JTBD):

"When my repeat task hits the third time, I want Claude to recognize the pattern and help me decide what to do — without forcing me to flatten what makes my work mine."

Five hypotheses tested. Three confirmed. One reframed. One inverted

Hyp.	Statement	Verdict	Evidence
H1	Re-explaining context is real friction for Power Repeaters	Confirmed 	80% frustrated overall · 90% in Power Repeater cohort · “frustrating and time consuming.”
H2	The third-task moment is the right nudge trigger	Confirmed 	52% prefer nudges after several repeats · 96% accept 5-min-for-2-min tradeoff
H3	Setup friction is the main conversion barrier	Reframed 	Top barrier is “I don’t repeat enough” (32%) · setup confusion only 12%
H4	Output quality fear blocks Skill adoption	Confirmed 	76% fear worse performance · 40% “very concerned”
H5	Consistency matters more than time savings	Inverted 	Time savings drove positive response · 96% accepted Q14



Survey Q10 · Multi-select · n=25

*Honest research changes the product.
We let it.*

*n=25 survey · 5 in-depth interviews (Apr 2026) ·
Hypothesis log timestamped pre-research · Survey
Results · Interview Synthesis · Hypothesis Log*

3 things this evidence reshaped:

- 1 Pattern blindness > setup friction → nudge must reveal repetition users can't see themselves
- 2 76% fear quality drift → Dry-Run Simulator becomes mandatory, not optional
- 3 Time, not consistency → tradeoff card leads with minutes saved per future run

We considered three paths. Only one earns the third-task moment.

Each option could solve the 80% frustration. Only one survives the H4 quality-fear test and the H5 time-not-consistency reframe.

Saved Prompts Library

Manual · Low-effort

Concept A Notion-for-prompts inside Claude. User saves and tags reusable prompts manually. Searchable library on the side panel.

Strengths

- Familiar UX pattern
- Lowest engineering effort
- No surprise factor — user stays in control

Why it fails

This is the workaround we already proved doesn't work. 20% of surveyed users save prompts in Word/Notes — and they're still in the 80% frustrated cohort. Building a better filing cabinet doesn't fix pattern blindness.

Auto-Skill Creator

Aggressive · Zero-friction

Concept Detects 2nd repeat. Auto-creates a Skill in the background. User finds it in their library next session.

Strengths

- Zero user friction
- Maximum coverage of repeat tasks
- Highest theoretical reach

Why it fails

Background auto-saving breaks trust before value is proven. Microsoft Recall was branded a “privacy nightmare” in 2024 and rebuilt as opt-in after backlash. Aniruddha needs to review what’s saved before committing; auto-creation removes that choice.

SkillSense

Decision Layer · Confidence-first



Concept Detects 3rd same-shape repeat. Surfaces time math + dry-run preview. User decides in 60 seconds.

Strengths

- Honors H2 (third-task moment, 96% accept tradeoff)
- H4 dry-run unlocks quality fear
- H5 time-led framing earns the click
- Respects Reluctant Adopter veto

Why it wins

The only option that solves pattern blindness and preserves trust. Effort is highest, but the RICE math still wins because Confidence and Impact dominate.

“Option Auto-Skill Creator reaches more users on paper. We rejected it because silent auto-saving is exactly what got Microsoft Recall called a 'privacy nightmare' in 2024 — Microsoft delayed launch by a year to fix it. We're not making that mistake.”

The market builds execution tools. SkillSense is the decision before the tool. That's the only slot still open.

RICE Prioritisation · Scale 1–10

Solution	Reach	Impact	Confidence	Effort	RICE	Verdict
SkillSense	9	9	8	7	92	P1 — Build
Auto-Skill Creator	9	6	4	5	43	Reject
Saved Prompts Library	7	3	9	9	21	Reject

Survey

Interview Synthesis

RICE = (Reach × Impact × Confidence) ÷ Effort × 10. Higher = stronger. Full reasoning in linked Sheet.



The third time, SkillSense speaks. Not before. Not after.

Built around the only nudge moment users actually accept — H2 confirmed: 52% prefer nudges after several repeats, 96% accept the 5-min-for-2-min tradeoff.

Mon 9:14 AM



Write the weekly engineering update.



Mon (next week) 9:14 AM




Write the weekly engineering update (again).



Mon (week after) 9:14 AM



Write the weekly engineering update.



I've seen this shape 3 times in 14 days.
Want to save it as a Skill?



Pattern detected:
weekly eng update · 4 sections · same tone



Time math:
 $12 \text{ min} \times \sim 8 \text{ future runs} = 96 \text{ min saved/quarter}$

Dry-run preview (Tappable)



Save Skill (60s setup)

Not now

Don't ask for this one

Detect:

after 3rd same-shape repeat
(H1, H2). Respect the
threshold.

Risk Mitigation:

Dry-run runs in shadow
before any save. Show the
value before the ask. (H4 —
quality fear).

Decide:

60s commit or one-tap
dismiss. Time matters more
than consistency. (H5 —
time matters more than
consistency)

Two repeats are coincidence. Three is a pattern.

Survey

Interview Synthesis


*You decide once. Not 50 times. The product's only
job is to earn that one decision*

Every step earned by a confirmed hypothesis.

Six features identified. Three ship in V1. One we'll never build.

Every feature traces back to a hypothesis or a real interview voice. The Won't-Have is the line we won't cross.

THE V1 LOOP HAS THREE STAGES



"Three stages. Three Must-Have features. Each maps to a stage."

What V1 actually solves

Detect what Kishan can't see himself — 80% repeat context, 72% repeat tasks weekly

Preview what would break Vijay's skepticism — 76% fear a Skill performs worse than a fresh prompt

Decide in 60 seconds — Divyansh's frustration math, 96% accept the time tradeoff

What V1 does NOT solve:

- ✗ Editing saved Skills (V2 — Aniruddha's variation need)
- ✗ Cross-project sync (V3 — Aniruddha's per-project flavor)
- ✗ Public sharing (never — Vijay's role makes this a non-starter)

"A decision layer that judges reuse before any save. SkillSense doesn't build your Skills — it tells you which ones are worth building."

A decision layer that judges reuse before any save. SkillSense doesn't build your Skills — it tells you which ones are worth building.

FEATURE PRIORITISATION (MoSCoW)

Bucket	Feature	Why this verdict
Must	Pattern Detection Engine	Kishan: "Just detect the repeat for me and let me save it fast." 80% of users re-explain context they already typed. We're building exactly the detection he asked for.
Must	Dry-Run Simulator	Vijay: "I don't think I require it for now." That's the skepticism we have to break. 76% fear a saved Skill will perform worse — preview-before-commit shows value before the ask.
Must	Time-Saved Tradeoff Card	Divyansh: "frustrating and time consuming." Time saved drove 96% acceptance, not consistency. Card leads with minutes saved per future run.
Should (V2)	Skill Edit & Refinement	Aniruddha: "Show me the saved version preserves variation before I commit." V2 makes saved Skills editable, not rigid templates.
Could (V3)	Cross-project Skill Sync	Aniruddha again: "har project ka apna flavour hota hai" — every project has its own flavor. Sync must respect per-project variation.
Won't	Public Skill Marketplace	Two reasons. Vijay is already at "I don't require it" with private Skills — public publishing compounds the trust deficit. And as a Cybercrime Analyst, publishing his work patterns publicly is a non-starter.

Each feature scored against the JTBD: "When my repeat task hits the third time, I want Claude to recognize the pattern and help me decide what to do — without forcing me to flatten what makes my work mine."


Four interview voices. Four different verdicts. The product listens to all of them.



From keystroke to saved Skill. Sixty seconds. Zero guesses


USER FLOW · END-TO-END Both end-states are valid. Save earns the next nudge; dismiss earns silence. The product respects either choice.

[View on Figma](#)



★ **SUCCESS METRICS** - One North Star. Three leading metrics. One guardrail. Each tied to a survey number.


Type	Metric	Target	What it measures
★ NSM	Saved Skills used ≥ 3× in 30 days	40%+	Whether saves earn their place. Filters vanity saves and one-time curiosity by requiring 3 actual uses within a month
L1 Activation	Skill-save acceptance rate	≥40%	When the decision card fires, what % of users save. Validates the 3-of-3 threshold (H2 — third-task moment).
L1 Quality	Post-save quality CSAT	≥80%	After 3rd Skill use, user rates output vs fresh prompt. Catches quality drift (H4 — fear unlocked).
L1 Retention	30-day Skill-active retention	≥70%	Users with at least 1 Skill used 3+ times in 30 days. Proves time-saved framing earned the click (H5 inverted).
Guardrail	"Don't ask for this one" rate	≤15%	Pattern-mute frequency. Too high = nudging users who don't want nudging (Reluctant Adopter veto, quantified).



Five screens. Each earns one interview voice's trust.


Not every screen ships. These five decide if SkillSense does

FIRST LOAD · EMPTY STATE - Claude top-left. SkillSense earns its name only in the sidebar where it does the work. The 2-of-3 indicator surfaces the rule before it triggers — never on first try.




Sets the contract Vijay can read before deciding

DECISION CARD · PATTERN A - Inline in chat — never modal. Pattern Detected + Time Math split into two columns. Four CTAs in honest hierarchy: Dry-run preview as primary, Save as secondary, Not now / Don't ask as tertiary text links.




Kishan's request, in pixels: 'Just detect the repeat for me and let me save it fast.

DRY-RUN PREVIEW · GROQ-GENERATED - Real Llama 3.3 70B output, not mocked. Pattern-specific prompt generates pattern-appropriate sample. User sees the Skill's value before committing.




Vijay's skepticism unlock: 'I don't think I require it' becomes 'I can see why now.

SAVED · TIME SAVED · CONTROLS- Skill in sidebar with reversible controls (rename, delete with undo). Time Saved updates from 0 → ~96 min immediately. Three-dot menu = save isn't a one-way door.



Divyansh's frustration math, paid back: 96 minutes per quarter, visible every time he opens the app.

MULTI-PATTERN · PATTERN B- Three patterns wired in V1: weekly update, meeting summary, code review. Each detects independently, each fires its own card with pattern-appropriate time math.



Aniruddha's variation need respected: same engine, different patterns, no flattening.

Five screens. Four interview voices. Three confirmed hypotheses. One product, running.

 LIVE PROTOTYPE

Free at the moment of pain. Honest about what could kill us.

No new business model. No paid acquisition. Three risks named — each tied to the research that surfaced them.

How it makes money

Bundled in Claude Pro. Not standalone.

Free for Free users. Pattern detection plus 3 saved Skills max. Enough to feel the value, capped to drive Pro.

Bundled in Pro (\$20/mo). Unlimited Skills, full dry-run preview, three-pattern detection. Zero net new pricing — SkillSense is a Pro feature, not a separate product.

The honest argument: Anthropic's growth lever is not another SKU. It is retention. SkillSense reduces re-prompting waste and increases stickiness because saved Skills become switching costs.

How it reaches users

26M MAU already inside Claude. No paid acquisition.

Distribution we already own. 18.9M web + 7.38M mobile users. SkillSense ships natively to Pro — no install, no signup, no behavior change.

First 1,000 users: the Power Repeater cohort surfaces the feature organically. The 2-of-3 indicator does the discovery work — users see it the second time they repeat, then save the third.

Y1 SOM target: 720K active users. That is 5% of a 14.4M Power Repeater SAM, sized from Anthropic public data plus survey-derived behavioral filters.

“We don't need a new business model. We need to earn one decision at a time.”

“We're not acquiring users. They're already here. We're earning the third decision they make.”

What could kill it - Three risks. Each tied to research. Each mitigated in V1.

Detection threshold wrong

~**Risk:** Fires too early (annoys) or too late (useless).

~**Why this matters:** Setup confusion was only 12%, but if we mis-time the nudge, the H2 confirmation collapses.

~**Mitigation in V1:** Pattern threshold (3 repeats / 14 days) is configurable per pattern. Beta cohort validates before broad ship.

Quality drift in saved Skills

~**Risk:** Saved Skills produce worse output than fresh prompts.

~**Why this matters:** 76% already fear this. One bad Skill kills trust permanently.

~**Mitigation in V1:** Dry-Run Simulator benchmarks preview quality before save, so low-confidence outputs do not get formalized.

Reluctant Adopter churn

~**Risk:** Users like Aniruddha opt out in the first week.

~**Why this matters:** Time saved earns the click, but dignity and variation determine whether users stay.

~**Mitigation in V1:** “Don't ask for this one” preserves choice. Pattern mute is tracked as signal, not failure, and dismissal stays reversible.